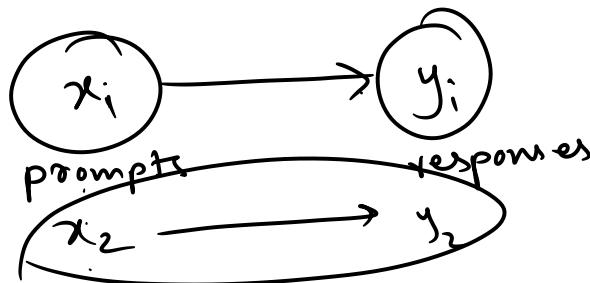


Fine-tuning

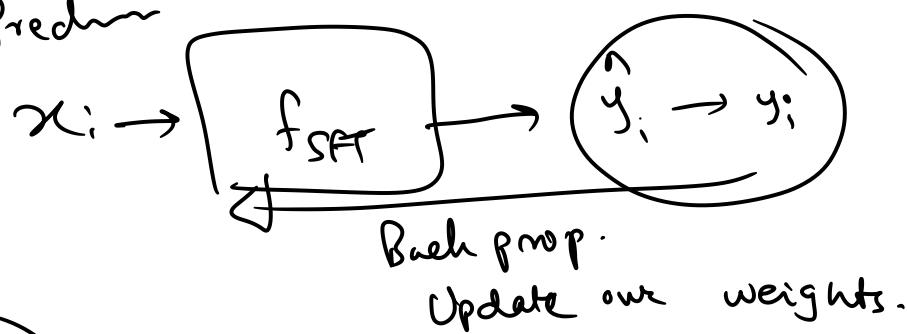
Supervised fine-tuning

$$f_{SFT} = f_{\text{pretrained}} =$$

Collect data



Next Word Prediction



f_{SFT}

Fine-tuning

SFT

Data Quality ↗

Pr → Response
Oh - f

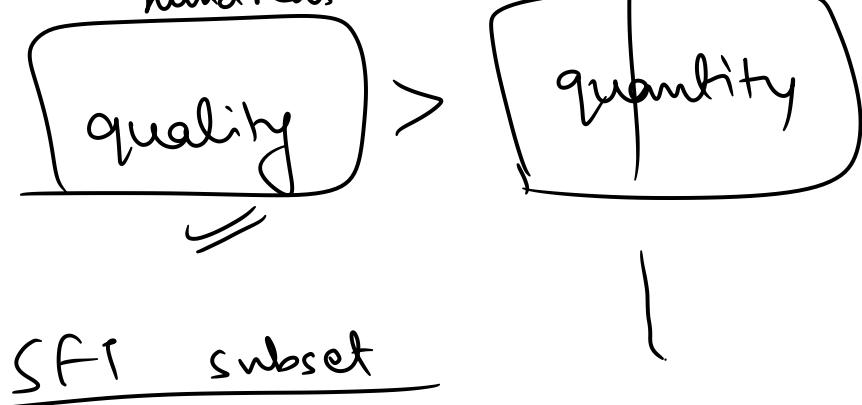
update: all the model weights & bias are updated.

Obs → pairs : 100-200

hundreds ↘

fine-tun.

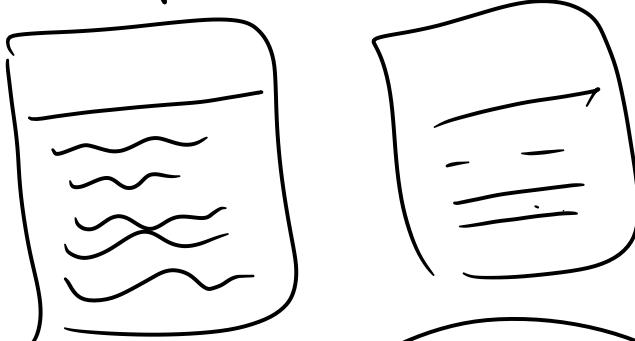
Mimic



Instruction Tuning:

It is a fine-tuning method that improves the ability of an LLM to perform on unseen data.

LLM
Poems



Challenges:

- 1) Very high quality data.
- 2) Very sensitive prompt dist.
- 3) Computationally expensive

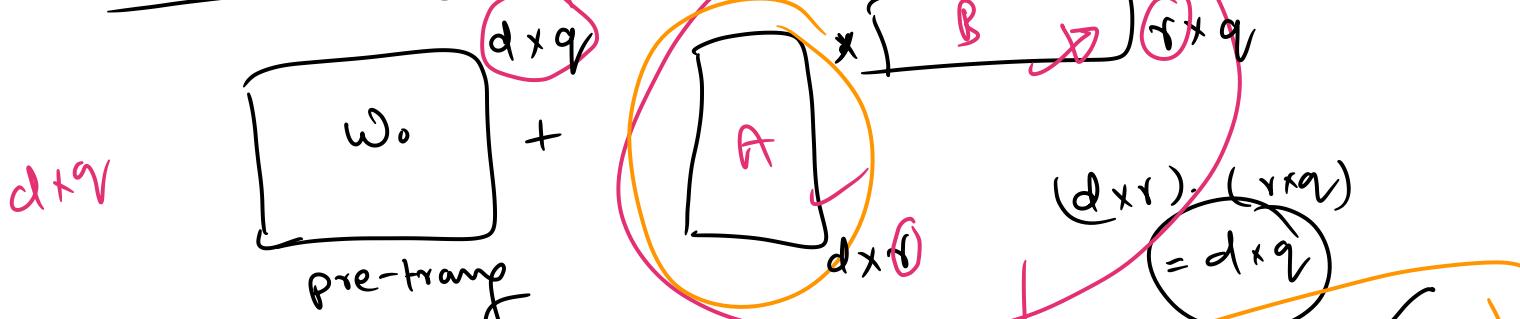
instruction tuned model
Param + Way Billions & trn

Each time updating all the weights
 w_1

Parameter Efficient Fine-tune

Updation of only a subset

LORA (low Rank Adaptation)



$$w = w_0 + AB$$

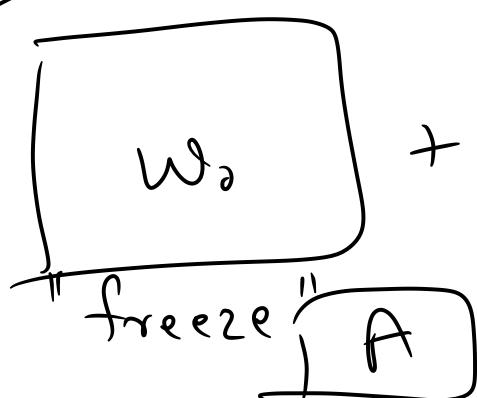
$$\delta \ll \min(d, q)$$

Rank $\leq \text{Rank}(A, B)$

d, q

MT H-113
Low Rank

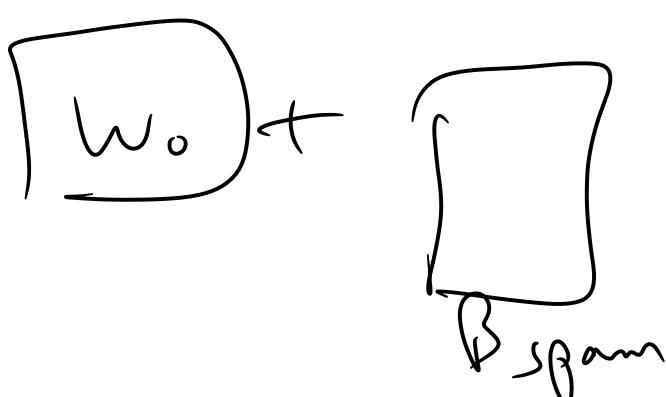
LORA



Traffic (fine-tune)

Fine-tune

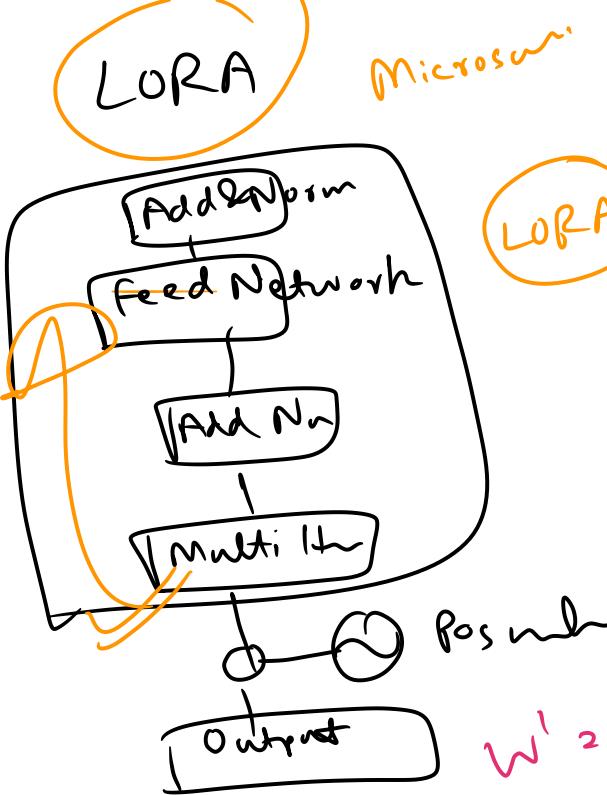
Spam Detection



A_{spam}

→ Spam

Spam
Detection

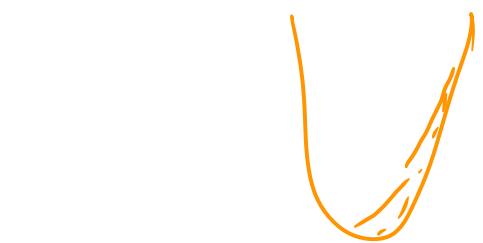


Experiment LORA features

LORA P =
 → LORA needs a higher learning rate than full fine-tune
 → LORA does poorly on large batch size compared to full fine-tune

$$w' = w - \eta \frac{\partial J}{\partial w}$$

Rank d

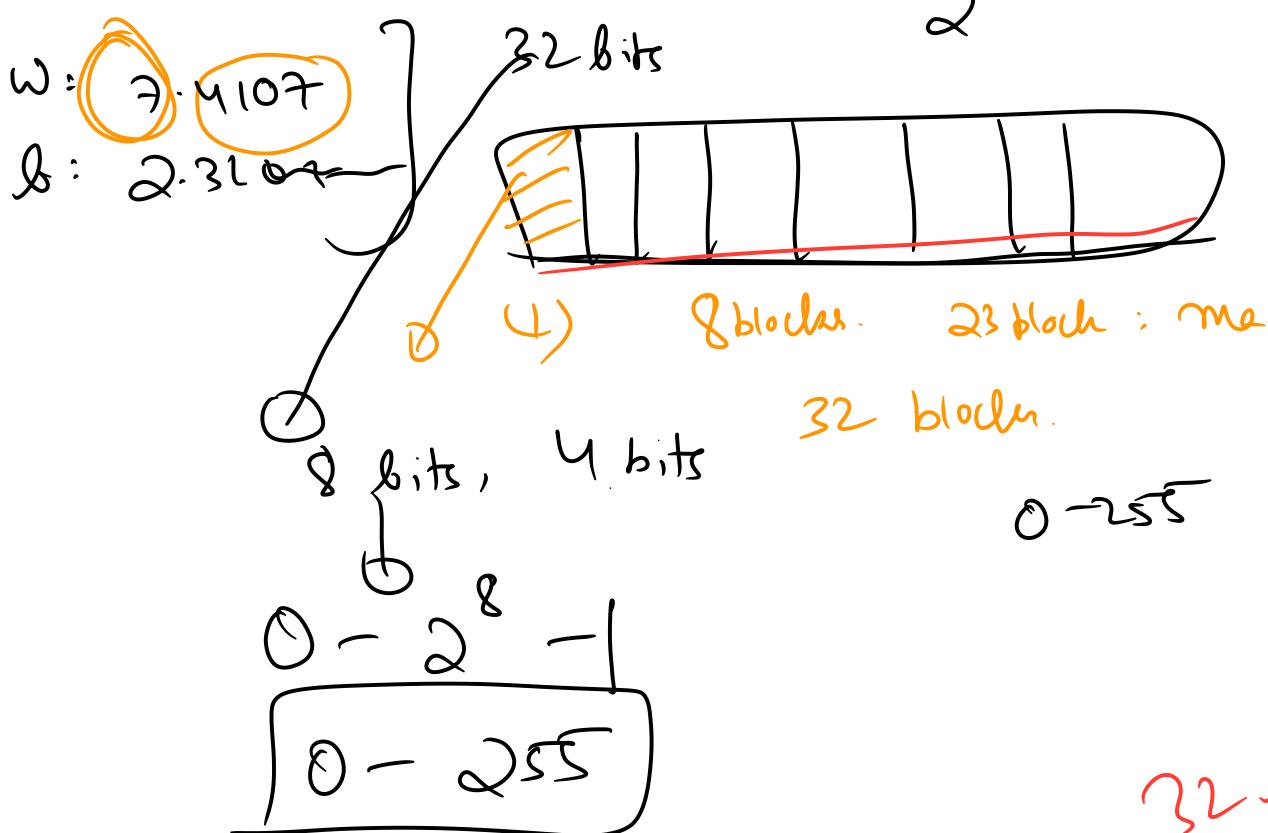


weights, bias



Computational Cost ↓
 Memory → Efficiently

Quantizator



Weight: $[10, 100, 150, \dots, 1000]$

~~Binary~~

$SF = \frac{1000}{255} = \frac{0}{0} = \frac{1000}{255} = \frac{0}{0}$

$\pm \frac{1000}{255} \Rightarrow \frac{0}{3} \rightarrow \frac{100}{3} = \dots$

$0 \rightarrow 1000$

32 bits

4 bits

min-max
3 - Non

$-30, -10, 0, -50, 100$



$$SF = \frac{100 + 30}{255} = \frac{130}{255} =$$

(compute hr)

$$\begin{matrix} -30 \\ 0 \\ 100 \end{matrix} = \dots + \alpha = 0$$

$\frac{-10}{SF} + \alpha, \frac{0}{SF} + \alpha,$



More Saturations
Loss of Info.

QLOFT

