

Assignment - 2

Date: / / Page No.:

1 A perceptron computes a linear decision boundary of the form -

$$y = \text{sign}(\omega_1 x_1 + \omega_2 x_2 + b)$$

thus it can separate data that is linearly separable
Now,

Truth table of XOR -

x_1	x_2	XOR
0	0	0
0	1	1
1	0	1
1	1	0

XOR is not linearly separable, and a single perceptron can only learn linear boundaries

Hence, a perceptron fails on XOR.

Now,

When we introduced multi-layer perceptrons, hidden layers with non-linear activation functions were added. These layers transform the input into a space where XOR becomes linearly separable, allowing the network to learn non-linear decision boundaries.

2 A linear layer computes

stack two linear layers:

this is still of form: $y = W_2(W_1x + b_1) + b_2 = (W_2W_1)x + (W_2b_1 + b_2)$

$$y = Wx + b$$

thus linear layers stacked on linear layers remain linear

- During backpropagation, gradients are computed using the chain rule:

$$\frac{\partial L}{\partial w} = \prod \text{activation derivatives}$$

if activation derivatives are less than 1 (eg - sigmoid)
then multiplying many small numbers \rightarrow exponentially smaller gradients

ex. Earlier layers receive almost zero updates

- Sigmoid activation:

$$\sigma'(x) = \sigma(x)(1 - \sigma(x)) \leq 0.25$$

if derivative is always < 1

ex. saturates at 0 or 1 \rightarrow gradients ≈ 0

Now, ReLU activation:

$$f(x) = \max(0, x), \quad g'(x) = \begin{cases} 1 & x > 0 \\ 0 & x \leq 0 \end{cases}$$

thus, derivative = 1 for active neuron and no saturation for positive inputs

Hence gradients flow better through deep networks with ReLU

3) Transformers use self-attention which is permutation-invariant. Without positional information, the model cannot distinguish: ex - "dog bites man" vs "man bites dog".

thus positional encoding is needed to inject order of tokens. And let attention know which token comes before or after which.

Aspect.	Sinusoidal PE	Absolute PE
1. Type	fixed, deterministic	Trainable parameters
2. Formula	Uses sin & cos of different frequencies	Lookup table of embeddings
3. Generalisation	Can extrapolate to longer sequences	Limited to trained length
4. Parameters	No extra parameters	Add parameters
5. Relative info	Implicit	Weak

RoPE rotates query and key vectors based on position:

$$\text{this gives: } q' = R(\theta_{\text{pos}})q, \quad k' = R(\theta_{\text{pos}})k.$$

↳ Relative positional encoding naturally

↳ Better extrapolation to long sequences

↳ Stable dot products

4 Query (Q): what the current token is looking for

Key (K): what each token offers

Value (V): the info. to be passed if a key matches the query

Attention computes:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

Now,

- Dot products $Q \cdot K$ grows in magnitude as dimension d_k increases
- Large values \rightarrow softmax becomes too sharp
- Gradients become very small

Scaling by $\sqrt{d_k}$:

$\frac{QK^T}{\sqrt{d_k}}$ keeps scores in a stable range, preventing saturation and improving training

- Now,
- ⇒ Diagonal entries correspond to a token attending to itself
 - ⇒ A token's query and key are derived from the same embedding
 - ⇒ So, similarity $Q_i \cdot K_i$ is usually maximal

5 Multi-head attention allows the model to look at the same sequence from different perspectives at the same time.

Each head can learn to focus on:

- syntax (e.g. subject - verb relations)
- semantics (meaning-related tokens)
- long-range vs short-range dependencies

This increases expressive power without increasing total computation too much

- d_{model} - Total embedding dimension of the model
- h - number of attention heads
- d_{head} - Dimension of each head

$$d_{\text{head}} = \frac{d_{\text{model}}}{h}$$

6 Greedy decoding selects the most probable token at each step:

$$y_t = \arg \max P(y_t | y_{<t}, x)$$

It makes locally optimal choices, ignores future consequences and cannot revise earlier decisions

So it may miss the globally most probable sequence

- Beam Search keeps the top-K (beam width) partial sequences, Expands all candidates at each step. Also, selects the sequences with highest total log-probability

example where greedy decoding fails

Assume probabilities

Time step 1

- "I" : 0.6
 - "You" : 0.4
- Greedy picks "I"

Time step 2

After "I" → "am" : 0.4

After "I" → "like" : 0.6

Greedy picks "like"

$$\text{Sequence prob: } 0.6 \times 0.6 = 0.36$$

But

"You" → "are" : 0.9

$$\text{Sequence prob: } 0.4 \times 0.9 = 0.36 \text{ (tie here)}$$

more clearly:

"I" → "like": $0.6 \times 0.5 = 0.30$

"You" → "are": $0.4 \times 0.9 = 0.36$

Beam search ($k=2$) keeps both prefixes and
selects "You are", which ~~already~~ misses